

Founder Approval and Authority Status

Field	Controlling Value
Founder Approval Status	**FOUNDER APPROVED**
Founder Approval Date	**July 16, 2026**
Constitutional Status	**CONTROLLING SUBSTANTIVE CANON**
Canon Adoption	**PENDING SEPARATE ADOPTION EVIDENCE**
Canon Lock	**NOT LOCKED BY THIS DIRECTIVE**
Implementation Authority	**FALSE**
Schema / Migration Authority	**FALSE**
Runtime Activation Authority	**FALSE**
Production Authority	**FALSE**
External-Provider Activation Authority	**FALSE**
Public-Claims Authority	**FALSE**
Public-Launch Authority	**FALSE**

Founder approval establishes the controlling substantive constitutional text. Adoption, lock, implementation, runtime, production, provider, public-claims, and launch authority remain separately gated.

MASTER CONFIGURATION AND

FEATURE FLAG GOVERNANCE MODEL

Version 2.0

Document Classification	Constitutional Governance Canon
Founder Decision Status	FD01-FD50 APPROVED AS RECOMMENDED
Drafting Status	Controlled Canon Draft
Canon Adoption	NOT AUTHORIZED BY THIS DOCUMENT
Canon Lock	NOT AUTHORIZED BY THIS DOCUMENT
Implementation Authority	FALSE
Runtime / Production Authority	FALSE
Public Launch Authority	FALSE
Prepared For	EquineSync
Effective Date	Upon separate founder adoption

Controlling Principle

Configuration may change how EquineSync behaves operationally only within expressly delegated boundaries. It may never change, weaken, bypass, suspend, or redefine what EquineSync is constitutionally obligated to protect.

Document Control and Navigation

This document follows EquineSync's approved constitutional canon structure. Section numbering, control status, authority boundaries, cross-canon dependencies, founder decisions, implementation limits, and adoption status are stated expressly.

Table of Contents

1. Purpose and Constitutional Function
2. Scope and Applicability
3. Definitions
4. Foundational Principles
5. Governance Authority and Decision Rights
6. Configuration Classification and Risk Tiers

7. Configuration Scope, Precedence, and Inheritance
8. Immutable and Nonconfigurable Controls
9. Feature Flag Categories and Lifecycle
10. Approval, Dual Control, and Change Evidence
11. Rollout, Experimentation, and Validation
12. Emergency Controls and Safe Degradation
13. Horse Welfare and Safeguarding Boundaries
14. AI Configuration Governance
15. Financial Configuration Governance
16. Jurisdictional and Regulatory Configuration
17. Offline, Synchronization, and Device Configuration
18. External Provider and Integration Configuration
19. Metadata, Registry, Versioning, and Retirement
20. Audit, Monitoring, Enforcement, and Exceptions
21. Cross-Canon Dependencies and Conflict Resolution
22. Founder Decision Ledger FD01-FD50
23. Implementation Boundary
24. Adoption, Lock, and Amendment Status

1. Purpose and Constitutional Function

This Master Configuration and Feature Flag Governance Model establishes the controlling constitutional governance for runtime configuration, operational settings, entitlement controls, feature flags, release controls, emergency disablement mechanisms, and related platform behavior across EquineSync.

Its purpose is to ensure that configuration remains a bounded operational mechanism rather than a means of rewriting constitutional truth, expanding authority, weakening safeguards, or bypassing adoption, implementation, release, or production gates.

2. Scope and Applicability

This canon applies to every EquineSync product, environment, tenant, organization, facility, barn, user role, horse record, mobile or web client, offline mode, integration, provider, AI capability, workflow, and future deployment model.

It governs platform, jurisdictional, organization, facility, barn, role, user, horse-specific, session, device, provider, and emergency configuration, together with the full lifecycle of feature flags.

3. Definitions

Configuration means an approved, bounded value or setting that changes permitted operational behavior without altering constitutional obligations or creating new authority.

Feature flag means a governed runtime control used for staged release, pilot access, experimentation, migration, compatibility, provider transition, operational containment, or emergency disablement.

Immutable control means a constitutional or mandatory safeguard that may not be weakened, disabled, or redefined through ordinary configuration or feature flags.

Safe default means the documented behavior that applies when configuration is absent, invalid, expired, conflicting, unavailable, or not yet approved.

4. Foundational Principles

Constitutional canon controls over configuration. Configuration may refine permitted operations but may not amend governing truth.

Missing or invalid configuration never creates expanded authority. EquineSync must fail closed, preserve existing truth, or degrade safely according to the governing domain.

Configuration flexibility is subordinate to identity integrity, authorization, privacy, audit evidence, financial truth, safeguarding, horse welfare, legal obligations, and AI authority boundaries.

Configuration, feature flags, and emergency controls must be attributable, reviewable, reversible where safe, and supported by durable evidence.

5. Governance Authority and Decision Rights

New configuration may be proposed by Engineering, Product, Operations, Security, Privacy, Compliance, Finance, Customer Success, Implementation, Trust and Safety, designated equine-domain authorities, or another founder-approved body.

Proposal authority does not create approval, implementation, deployment, runtime, or production authority.

Material configuration requires product ownership, technical architecture, security, affected-domain governance, and release approval. Founder approval is required for new constitutional-impacting configuration classes, expanded override authority, or any conversion of an immutable rule into configurable behavior.

6. Configuration Classification and Risk Tiers

Class	Name	Examples	Minimum Governance
C0	Personal Preference	Display, accessibility, personal notification, and nonconsequential user preferences	Single authorized user; standard audit where appropriate
C1	Routine Operational	Bounded workflow, template, and routine tenant settings	Authorized role; ordinary validation and audit
C2	Material Operational	Cross-user, cross-workflow, integration, or authoritative-record settings	Documented justification, testing, rollback, and elevated approval
C3	High Risk	Security, privacy, AI, finance, safeguarding, welfare, exports, retention, provider, or jurisdiction settings	Dual approval, enhanced evidence, monitoring, rollback, and periodic review
C4	Constitutional or Prohibited	Controls that may not be changed through ordinary configuration	Constitutional amendment or separate founder governance required

7. Configuration Scope, Precedence, and Inheritance

The default precedence order is: constitutional canon; jurisdictional and regulatory requirements; platform-wide approved configuration; organization; facility; barn or operational unit; role or user; horse-specific; and session, device, or temporary local preference.

Lower levels may override higher levels only where the higher-level setting expressly permits override, the actor has authority, the value remains within the permitted range, dependencies remain valid, and the result does not weaken a constitutional safeguard.

The most restrictive valid requirement controls unless a controlling canon expressly establishes a different rule.

8. Immutable and Nonconfigurable Controls

Ordinary configuration and feature flags may not weaken identity, authentication, authorization floors, audit generation, evidence preservation, privacy rights, encryption, key management, ownership truth, transfer continuity, financial integrity, safeguarding, minor protections, legal retention minimums, AI authority limitations, horse welfare minimums, legal holds, or mandatory regulatory controls.

A user-facing capability may be hidden or degraded only where the underlying constitutional protection remains active and an approved continuity path remains available.

9. Feature Flag Categories and Lifecycle

Approved categories are Release, Operational, Emergency, Migration, Provider Transition, Experiment, Internal Test, Customer Pilot, Jurisdictional Availability, Compatibility, and Degraded-Mode flags.

Every flag must have an owner, purpose, scope, risk class, creation date, intended removal date, maximum duration, review checkpoint, validation criteria, rollback plan, monitoring plan, and retirement criteria.

Permanent feature flags are prohibited. Long-lived variation must be converted into governed configuration, entitlement, capability registration, or jurisdictional policy. The default maximum life of a temporary flag is 180 days unless a shorter period is required by risk.

10. Approval, Dual Control, and Change Evidence

Dual approval is required for material changes affecting security, identity, authorization, privacy, safeguarding, finance, payments, exports, retention, legal holds, AI, integrations, encryption, emergency controls, jurisdictional compliance, platform defaults, or high-risk horse-welfare functionality.

Every material change must record business purpose, requesting actor, approving actors, scope, prior value, new value, effective time, justification, dependencies, validation, monitoring, rollback, and governing authority.

Backdating is prohibited. Prospective scheduling is permitted only with explicit timezone, approval state, conditions, cancellation authority, and conflict detection.

11. Rollout, Experimentation, and Validation

The default rollout path is development or synthetic testing; internal validation; controlled nonproduction testing; designated pilot; limited customer rollout; expanded cohort rollout; regional or jurisdictional rollout where applicable; general availability; and flag retirement with code cleanup.

A/B testing is limited to low-risk, nonconsequential matters such as layout, wording, navigation, discoverability, onboarding sequence, and optional productivity tools.

Experimentation is prohibited where variation could affect horse welfare, medical or treatment truth, medication administration, emergency response, identity, authorization, ownership, custody, transfer, safeguarding, minors, privacy rights, consent, audit evidence, claims, retention, payments, accounting, tax, reconciliation, deletion, exports, legal notices, or consequential AI decisions.

12. Emergency Controls and Safe Degradation

Emergency controls must exist where appropriate for integrations, AI providers, outbound notifications, synchronization, imports, exports, payment initiation, provider APIs, newly released features, background jobs, webhooks, automated communications, migrations, and destructive processes.

Emergency activation is limited to designated incident-response personnel and must capture actor, authority, timestamp, environment, scope, reason, incident reference, prior state, resulting state, expected impact, communications, validation, restoration criteria, restoration evidence, and post-incident review.

Emergency operation must favor safe degradation over uncontrolled continuation.

13. Horse Welfare and Safeguarding Boundaries

Mandatory horse-welfare protections may never be fully disabled. Configuration may alter presentation, local terminology, workflow order, or noncritical reminder cadence only where minimum safety and continuity remain intact.

Emergency contacts, authorized medication and treatment visibility, critical allergies and contraindications, feeding and water-care instructions, critical care schedules, horse identity, responsible-party information, emergency movement or custody information, incident documentation, and care-history preservation must remain available through an approved continuity path.

Customer customization may not reduce safeguarding, minor protections, or constitutional welfare minimums.

14. AI Configuration Governance

Customers may configure approved optional AI capabilities, drafting assistance, summarization preferences, tone, language, notification timing, module availability, and confidence thresholds within governed ranges.

Human review, permission inheritance, provenance, uncertainty disclosure, fact-versus-inference distinction, refusal when evidence is insufficient, privacy protection, minor safeguards, prohibition on unauthorized cross-tenant learning, provider training restrictions, non-determinative authority, and the prohibition on silent authoritative mutation are immutable.

Customers may disable optional AI globally or by organization, facility, module, approved role, or workflow. Disabling AI may not disable security, privacy, audit, governance, or non-AI continuity paths.

15. Financial Configuration Governance

Permitted financial settings may include currencies, tax mappings, approved fee schedules, invoice presentation, due-date rules, lawful late-fee settings, payment-method availability, approved processor selection, deposit settings, refund workflows, accounting mappings, approval thresholds, and notifications.

Configuration may not alter historical truth, ledger integrity, provenance, reconciliation, original transaction values, payment-state truth, refund evidence, dispute evidence, duplicate-payment protection, or required retention. Configurable financial values must be versioned by effective date.

16. Jurisdictional and Regulatory Configuration

Jurisdictional variation must be implemented through approved policy packages identifying jurisdiction, authority, effective date, covered features, required defaults, prohibited behavior, retention rules, notice requirements, data-location requirements, tax or payment requirements, minor protections, review date, and supersession history.

Capabilities requiring legal, contractual, processor, or regulatory approval remain disabled until required evidence exists. Absence of a blocking value is not permission to enable.

Mandatory legal, privacy, safeguarding, security, financial, audit, retention, welfare, and regulatory requirements cannot be opted out of by customers.

17. Offline, Synchronization, and Device Configuration

Offline-capable clients must identify the configuration version in use, detect stale or conflicting configuration, prevent local authority expansion, preserve integrity, reconcile upon reconnection, and generate audit evidence for material offline behavior.

When a device cannot confirm current high-risk configuration, it must fail closed or enter an approved degraded mode. Offline operation may not silently weaken constitutional protections.

18. External Provider and Integration Configuration

Provider configuration must remain provider-neutral where required, separate provider selection from provider activation, preserve minimum-necessary disclosure, identify shared data categories, enforce retention and region constraints, support revocation, rollback, vendor exit, and prevent fallback to an unapproved provider.

Configuration may reference a secret or credential identifier but may not expose secret values. Configuration authority and secret-management authority must remain separate.

19. Metadata, Registry, Versioning, and Retirement

Every material configuration item and feature flag must have a unique identifier, human-readable name, description, purpose, category, owners, allowed values, data type, default and safe default, scope, precedence, inheritance and override rules, authority, approvals, risk class, dependencies, incompatibilities, affected modules and data, jurisdictions, canon references, lifecycle dates, rollback method, validation, monitoring, audit class, secret-handling class, and retirement criteria.

Material configuration must maintain immutable version history and effective intervals. Rollback may not erase evidence or restore a value that is no longer lawful, safe, compatible, or constitutionally permitted.

Flags must be reviewed at least quarterly and retired when rollout, pilot, migration, provider transition, emergency, or experiment concludes.

20. Audit, Monitoring, Enforcement, and Exceptions

All material configuration and feature-flag activity must be auditable. Sensitive values must be redacted or represented by secure references or fingerprints rather than written in plaintext logs.

Configuration that violates constitutional governance must be rejected. Invalid, expired, absent, corrupted, conflicting, or unsynchronized high-risk configuration must trigger safe default behavior.

Exceptions require explicit governing authority, narrow scope, defined duration, evidence, monitoring, and review. No exception may silently amend constitutional canon.

21. Cross-Canon Dependencies and Conflict Resolution

This canon must be interpreted consistently with EquineSync's Identity, Agreement and Authorization, Privacy, AI, Audit, Claims, Record Stewardship, Communication and Notice, Financial Truth, External Architecture, Platform Operations and Release, Encryption and Key Management, Resilience and Recovery, Relationship, Horse Transfer and Continuity, Offline and Synchronization, Safeguarding, Minor Protection, and Horse-Welfare governance.

Where conflict is identified, the higher-order controlling canon must be preserved, the conflict recorded, silent resolution avoided, and founder review requested where necessary. The affected configuration remains nonoperative until the conflict is resolved.

22. Founder Decision Ledger FD01-FD50

Decision	Founder-Approved Rule
FD01	Configuration includes approved runtime settings, workflow options, UI behavior, notification preferences, integration parameters, operational defaults, and business parameters, but excludes constitutional truth and legal obligations.
FD02	Ownership, identity, authorization, audit, privacy, safeguarding, financial truth, legal retention, AI boundaries, and horse-welfare protections are not ordinary configuration.
FD03	Engineering, Product, Operations, Security, Privacy, Compliance, Finance, Customer Success, Implementation, Trust and Safety, and approved equine-domain authorities may propose configuration; implementation remains separately authorized.
FD04	New configurable behavior requires risk-proportionate Product, Architecture, Security, Domain Governance, and Release approval, with founder approval for constitutional impact.
FD05	Precedence is Constitutional Canon, Jurisdiction, Platform, Organization, Facility, Barn, User or Role, Horse, then Session or Device preference.
FD06	Lower-level override is allowed only when expressly permitted and may not weaken a higher-order safeguard.
FD07	Security, privacy, audit, legal, identity, AI safety, finance, safeguarding, and welfare rules inherit automatically.
FD08	Identity, authentication, authorization floors, audit, finance, ownership, safeguarding, retention, encryption, AI authority, welfare, privacy, and mandatory compliance remain locked.
FD09	Feature flags are governed runtime controls for release, pilot, testing, migration, provider transition, compatibility, containment, and emergency disablement.
FD10	Minimum identity, audit, financial, ownership, welfare, safeguarding, privacy, legal-hold, and retention protections may not be disabled.
FD11	Approved flag categories are Release, Operational, Emergency, Migration, Provider Transition, Experiment, Internal Test, Customer Pilot,

	Jurisdictional Availability, Compatibility, and Degraded Mode.
FD12	Permanent feature flags are prohibited; long-lived variation must become governed configuration or entitlement.
FD13	Temporary flags default to a maximum 180-day lifetime with review and documented extension.
FD14	Mandatory horse-welfare protections may never be fully disabled.
FD15	Critical welfare and emergency continuity information must remain available during degradation or rollback.
FD16	Customer safety customization may alter presentation and noncritical workflow only, not constitutional minimums.
FD17	Optional AI assistance, drafting, summaries, tone, language, timing, module availability, and thresholds within governed ranges may be configured.
FD18	AI human review, permission inheritance, provenance, uncertainty disclosure, privacy, safeguarding, non-determinism, and non-silent mutation remain immutable.
FD19	Customers may disable optional AI without disabling governance, security, privacy, audit, or non-AI continuity.
FD20	Configuration authority is role-scoped across platform, security, organization, facility, barn, professional, owner, guardian, and user levels; titles alone do not grant authority.
FD21	Dual approval applies to high-risk configuration, with narrow emergency exceptions subject to retrospective review.
FD22	Material changes require documented purpose, scope, risk, dependencies, validation, rollback, review date, and governing authority.
FD23	Rollout follows development, internal validation, nonproduction, pilot, limited release, expanded cohorts, jurisdictional rollout, GA, and retirement.
FD24	Subscription tiers may unlock optional capabilities but may not reduce constitutional protections.
FD25	Jurisdictional rollout is permitted through approved policy packages.
FD26	Regulatory-dependent capabilities remain disabled until required approval evidence exists.
FD27	A/B testing is permitted only for low-risk, nonconsequential UX and productivity matters.
FD28	Experimentation is prohibited in consequential welfare, medical, emergency, identity, ownership, safeguarding, legal, privacy, audit, payment, retention, deletion, export, and AI domains.
FD29	Material experiments require notice where law, contract, privacy, ethics, or reasonable expectation requires it.
FD30	Emergency controls must exist for integrations, AI, notifications, synchronization, imports, exports, payments, providers, migrations, releases, and destructive processes where appropriate.
FD31	Emergency activation authority belongs only to designated incident responders within defined scope.
FD32	Emergency actions require immutable evidence of actor, reason, scope, state change, impact, validation, restoration, and review.
FD33	All material configuration changes must record before and after values, actor, authority, approvals, scope, timing, justification, and validation.
FD34	Backdating is prohibited; corrections must be prospective and explicitly linked to the error.
FD35	Prospective scheduling is permitted with timezone, approval, cancellation, rollback, and conflict controls.
FD36	Material configuration requires version history, effective intervals, rollback eligibility, and restoration evidence.
FD37	Tenants may create bounded configuration within approved schemas, but not arbitrary executable behavior or hidden authorization paths.
FD38	Constitutional controls remain globally governed; tenants may add protection but not reduce the global minimum.
FD39	Branding, terminology, workflows, notifications, calendars, lessons, boarding, layouts, task templates, care templates, approved services,

	integrations, and invoice presentation may vary within defined bounds.
FD40	Financial settings may include currencies, taxes, fees, invoice presentation, payment methods, approved processors, deposits, refunds, mappings, thresholds, and notices.
FD41	Ledger integrity, reconciliation, provenance, transaction history, payment-state truth, dispute evidence, duplicate protection, and financial retention remain immutable.
FD42	Regulatory variation must use approved jurisdictional policy packages.
FD43	Legal defaults may be updated only through qualified review, documented authority, impact analysis, notice planning, release approval, and evidence preservation.
FD44	Customers may not opt out of mandatory law, privacy, safeguarding, security, finance, audit, retention, welfare, or notice obligations.
FD45	Deprecation requires notice, replacement path, migration guidance, compatibility period, final removal date, preservation plan, and verification.
FD46	Every configuration domain requires accountable business, technical, governance, operational, and escalation ownership.
FD47	Every material configuration item requires complete governance metadata.
FD48	Flags must be reviewed quarterly and retired with code, test, documentation, dependency, audit, and tenant-cleanup confirmation.
FD49	Rollout success must consider technical health, integrity, authorization, audit, security, privacy, finance, support, adoption, accessibility, safeguarding, and horse-welfare impact.
FD50	Founder approval is required for new constitutional-impacting domains or flag categories, expanded override or activation authority, immutable-to-configurable conversion, customer-authored executable configuration, cross-tenant inheritance, automated AI configuration authority, or materially changed emergency authority.

23. Implementation Boundary

Founder approval of FD01-FD50 and preparation of this controlled canon draft do not authorize configuration mutation, feature-flag creation or activation, tenant-setting changes, schema mutation, migration, secrets modification, provider activation, AI activation, financial-rule changes, live experimentation, production access, deployment, runtime implementation, public trust claims, or public launch.

Any implementation, migration, activation, production, or public-release authority requires a separate founder directive and the applicable readiness, security, implementation, release, and production gates.

24. Adoption, Lock, and Amendment Status

Founder Decision Status: APPROVED AS RECOMMENDED.

Document Status: FOUNDER-APPROVED CONTROLLING SUBSTANTIVE CANON READY FOR FOUNDER REVIEW.

Canon Adoption: NOT AUTHORIZED BY THIS DOCUMENT.

Canon Lock: NOT AUTHORIZED BY THIS DOCUMENT.

Implementation Authority: FALSE.

Production Authority: FALSE.

Public Launch Authority: FALSE.

Amendment of constitutional boundaries or founder decisions requires a separate governed founder action.

Founder Disposition	APPROVED AS RECOMMENDED
Controlled Drafting Authority	AUTHORIZED
Implementation Authority	NOT AUTHORIZED
Production Authority	NOT AUTHORIZED